

# Encrypted Keyword Search Using Path ORAM on MirageOS

Rupert Horlick – [rh572@cam.ac.uk](mailto:rh572@cam.ac.uk)

June 7, 2016

# Introduction

- ▶ Final year undergraduate Computer Science student
- ▶ Undertook project over 9 months
- ▶ Implemented Path ORAM protocol, along with a file system and search module
- ▶ Wrote 10,000 word dissertation on the whole process

# Overview

Motivation

Solution

Implementation

Evaluation

Summary

# Overview

Motivation

Solution

Implementation

Evaluation

Summary

# Motivation

- ▶ Cloud storage's popularity demands a stronger emphasis on privacy
- ▶ Encryption hides data from cloud storage providers
  - ▶ But hinders the ability to search
- ▶ Homomorphic encryption makes encrypted search possible
  - ▶ But can leak up to 80% of queries!
- ▶ Can we have the best of both worlds?

# Overview

Motivation

**Solution**

Implementation

Evaluation

Summary

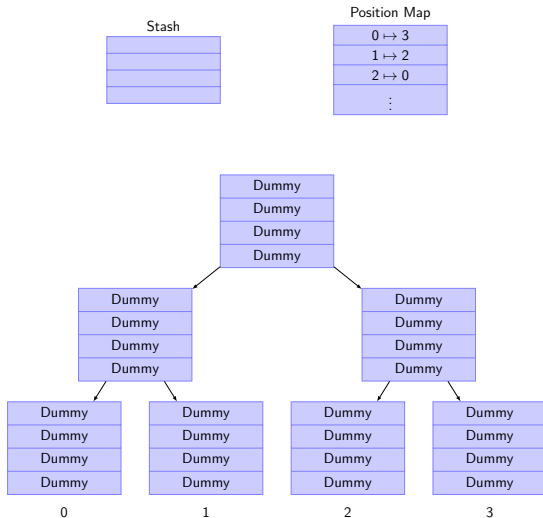
# Oblivious Random Access Memory (ORAM)

- ▶ A cryptographic protocol for obfuscating access patterns
  - ▶ Trusted client and untrusted storage server
  - ▶ Relies on cryptographically secure shuffling of data
- ▶ Originally applied to software protection
  - ▶ Repurposed for secure processors and cloud computing
- ▶ Original schemes had unacceptable overheads
  - ▶ Recent improvements have made ORAM more feasible

- ▶ Recent ORAM scheme (2013)
- ▶ Maintains three data structures
  - ▶ Binary tree on server
    - ▶ Each node is a bucket that contains up to  $Z$  blocks
    - ▶ Initially all blocks are dummy blocks
  - ▶ Stash on client
    - ▶ Working memory for blocks read from the tree
    - ▶ Initially empty
  - ▶ Position map on client
    - ▶ Associates to each block of data a leaf in the tree
    - ▶ Initially contains uniformly random values




# Path ORAM Initial Overview



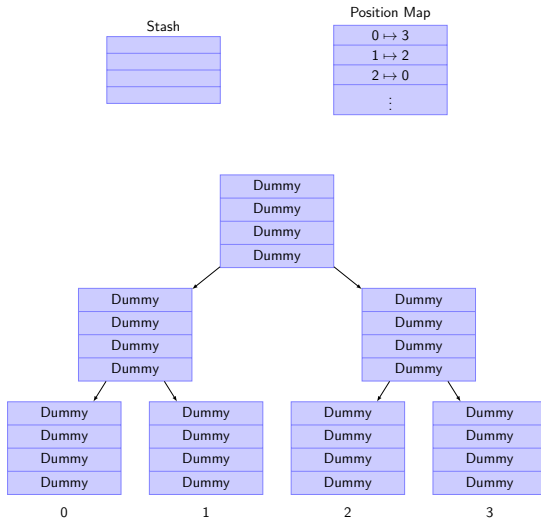
# Access Algorithm

- ▶ Take in parameters  $a$ , the address,  $op$ , the operation, and  $data^*$ , the new data for a write
- ▶ Then have the following steps:
  - ▶ Lookup position of  $a$  in position map,  $x$
  - ▶ Remap  $a$  to a random position
  - ▶ Read the  $x$ -th path into the stash
  - ▶ If  $op$  is write, then overwrite data for  $a$  with  $data^*$  in the stash
  - ▶ Write blocks from the stash back into  $x$ -th path
  - ▶ If  $op$  is a read, then return data

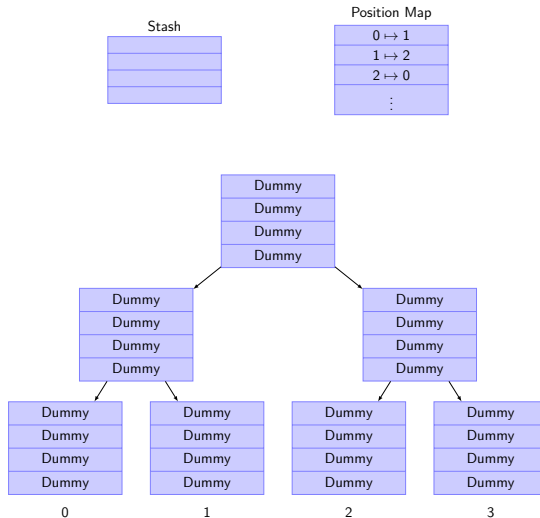
## Path ORAM Input




# Example Write: Lookup Position



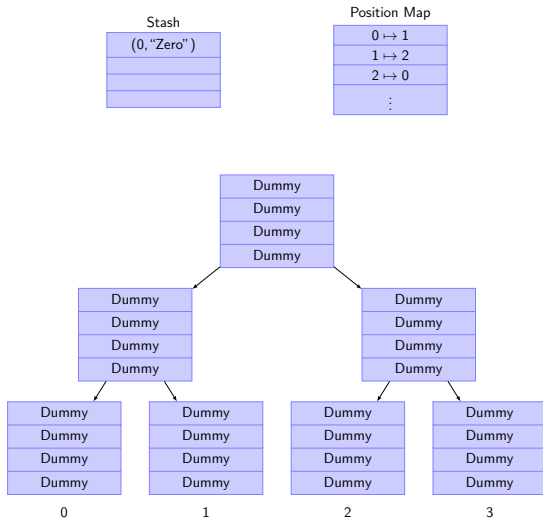
# Example Write: Remap Block



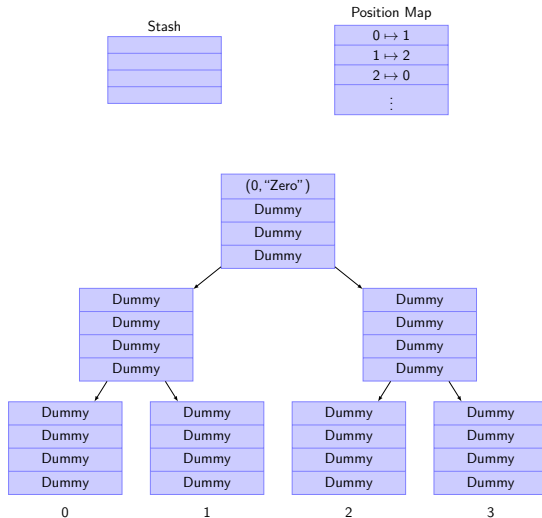
# Example Write: Read Path



# Example Write: Write Data




# Example Write: Write Path

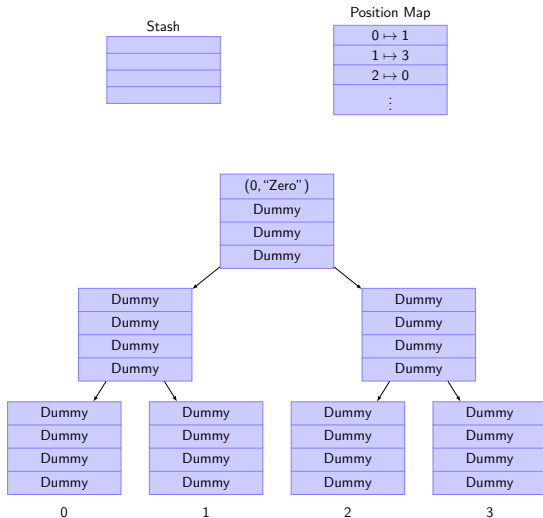





# Example Write: Lookup Position




# Example Write: Remap Block



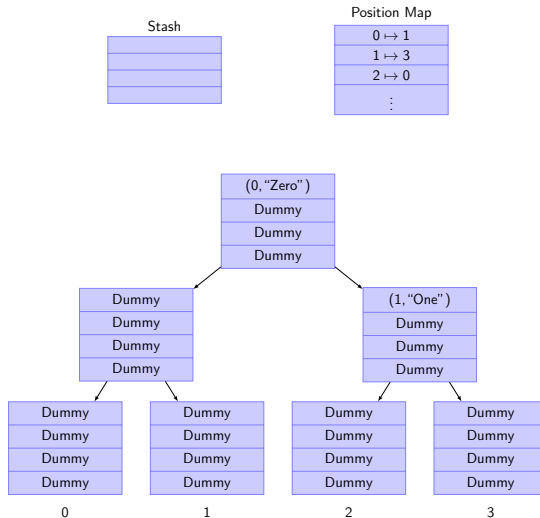
# Example Write: Read Path



# Example Write: Write Data



# Example Write: Write Path



# Introduction to MirageOS

- ▶ Library operating system, consisting of a command line tool and a set of libraries
- ▶ Compiles applications to a unikernel, a lightweight operating system
- ▶ Can be compiled for a number of targets including Xen and Unix
- ▶ Gives us control over where the ORAM application is deployed
- ▶ Allows a lightweight cloud instance to be spun up on demand

# Overview

Motivation


Solution

**Implementation**

Evaluation

Summary

# System Architecture



**Figure:** ORAM can be built as a MirageOS application, which can run on a trusted cloud instance.



# Implementation of Basic ORAM

- ▶ Built Path ORAM as a functor satisfying Mirage's BLOCK interface
- ▶ A functor is a module parameterised over the implementation of another module
- ▶ This allows the underlying implementation of the storage device to be abstracted away
- ▶ The stash was built as a hash table
- ▶ The position map was built as a 3-dimensional array to allow 64-bit addresses

# Recursive ORAM

- ▶ Do be able to disconnect from ORAM we need it to be stateless
- ▶ Writing the entire position map to disk is expensive
- ▶ Recursive ORAM stores the position map of the first ORAM in another ORAM
- ▶ This position map ORAM will have a smaller position map, because multiple addresses fit in one data block
- ▶ This process can be repeated until we have a constant sized position map
- ▶ Implemented this using recursive functors, parameterised in the position map

# The Rest

- ▶ On top of ORAM I built an inode-based file system
- ▶ This included an implementation of B-Trees to store the inode index
- ▶ On top of the file system I built a simple search module
- ▶ This consists of an inverted index built on a hash table of hash sets
- ▶ I incorporated an existing encryption library, also built as a functor

# Overview

Motivation

Solution

Implementation


**Evaluation**

Summary

# Parameter Optimisation

- ▶ ORAM has a number of parameters that can be adjusted
- ▶ These include the size of blocks and the number of blocks in a bucket
- ▶ I focused on block size, because this has been shown to dramatically affect the speed of IO operations
- ▶ The results showed this to be true, but there is a trade off to be made
- ▶ A large block size means a larger stash, so a larger overhead for statelessness

# Block Size Results

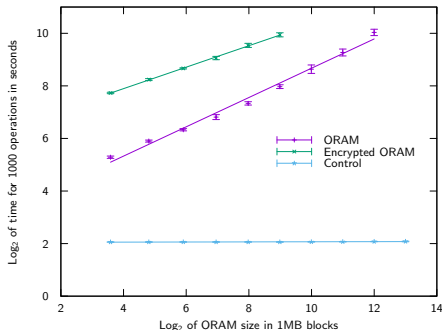


**Figure:** Plot of the time taken to transfer 80MB of data at varying block sizes and sizes of ORAM. Each line represents one ORAM size,  $N$ , so as block size increases, the time decreases.

# Performance Evaluation

- ▶ In the Path ORAM paper, they report a theoretical overhead of  $O(\log N)$
- ▶  $N$  here is the size of the ORAM in blocks
- ▶ I performed 10 runs of 1000 iterations, for a total of  $\approx 1\text{GB}$  per run
- ▶ I used ORAMs from size 12 (tree of depth 1) to 4092 (depth 9)
- ▶ The results showed a roughly logarithmic overhead, compared to the control
- ▶ Encryption added a further constant overhead

# Performance Results




**Figure:** The relationship between size of an ORAM in blocks and the time taken for 1000 operations, plotted for ORAM, encrypted ORAM, and a control block device with no ORAM. We take logs of both axes, because block size was increased in powers of two and we expect a log relationship.




# Security Evaluation

- ▶ Used statistical techniques to show that the observed access pattern is random
- ▶ The first technique is autocorrelation plotting
- ▶ This takes the correlation of a sequence with itself plotted for a number of lags
- ▶ For a random sequence, noise should cancel out to give values close to zero
- ▶ The second technique is runs testing
- ▶ This counts the number of runs of consecutive values all above or below the median
- ▶ We compare this number to that of a random process

# Autocorrelation Results




(a) Autocorrelation plot of a 200 iteration access pattern



(b) Autocorrelation plot of a 5,000 iteration access pattern

**Figure:** Two autocorrelation plots, with the autocorrelation coefficient on the y-axis and time lag on the x-axis. The dashed black lines represent confidence bands of 95% and 99%. For a random sequence, most of the points should fall within the 95% confidence bound, as they do on both of these plots.

# Runs Test Results



**Figure:** The distribution of the number of runs in 1000 access patterns of length 180. The dashed black lines represent 0.05% tail cut-offs. 92.2% of values fall within these bounds, implying that the access patterns were created from a random process.

# Overview

Motivation

Solution

Implementation

Evaluation

Summary

# Further Work

- ▶ Further exploration of the parameter space
- ▶ Examination of alternative performance measures including bandwidth overheads
- ▶ Deployment to the cloud and testing of overheads in real scenario

Thank you  
Questions?