

# Encrypted Keyword Search Using Path ORAM on MirageOS

Rupert Horlick – [rh572@cam.ac.uk](mailto:rh572@cam.ac.uk)

June 8, 2016

# Introduction

- ▶ Final year undergraduate Computer Science student
- ▶ Undertook project over 9 months
- ▶ Implemented Path ORAM protocol, along with a file system and search module
- ▶ Evaluated performance and security properties
- ▶ Wrote 10,000 word dissertation on the whole process

# Overview

Motivation

Solution

Implementation

Evaluation

Summary

# Overview

Motivation

Solution

Implementation

Evaluation

Summary

# Motivation

- ▶ Cloud storage's popularity demands a stronger emphasis on privacy
- ▶ Encryption hides data from cloud storage providers
  - ▶ But hinders the ability to search
- ▶ Homomorphic encryption makes encrypted search possible
  - ▶ But can leak up to 80% of queries! [Islam et al.]
- ▶ Can we have the best of both worlds?

# Overview

Motivation

**Solution**

Implementation

Evaluation

Summary

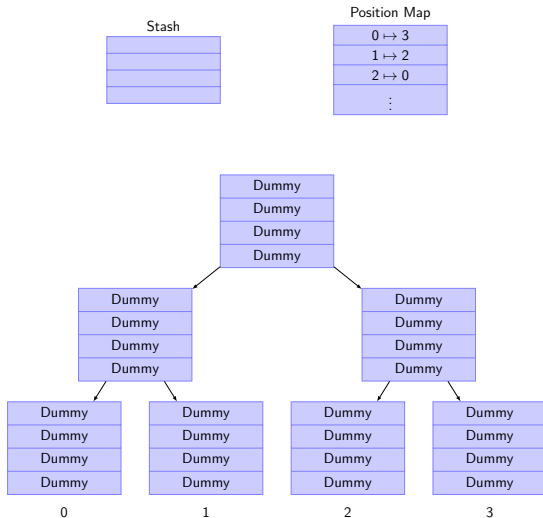
# Oblivious Random Access Memory (ORAM)

- ▶ A cryptographic protocol for obfuscating access patterns
  - ▶ Trusted client and untrusted storage server
  - ▶ Relies on cryptographically secure shuffling of data
- ▶ Originally applied to software protection
  - ▶ Repurposed for secure processors and cloud computing
- ▶ Original schemes had unacceptable overheads
  - ▶ Recent improvements have made ORAM more feasible

- ▶ Recent ORAM scheme (2013)
- ▶ Maintains three data structures
  - ▶ Binary tree on server
    - ▶ Each node is a bucket that contains up to  $Z$  blocks
    - ▶ Initially all blocks are dummy blocks
  - ▶ Stash on client
    - ▶ Working memory for blocks read from the tree
    - ▶ Initially empty
  - ▶ Position map on client
    - ▶ Associates to each block of data a leaf in the tree
    - ▶ Initially contains uniformly random values




# Path ORAM Initial Overview



# Access Algorithm

- ▶ Signature:  $\text{access}(a, \text{op}, \text{data}^*)$
- ▶ Then have the following steps:
  - ▶ Lookup position of  $a$  in position map,  $x$
  - ▶ Remap  $a$  to a random position
  - ▶ Read the  $x$ -th path into the stash
  - ▶ If  $\text{op}$  is write, then overwrite data for  $a$  with  $\text{data}^*$  in the stash
  - ▶ Write blocks from the stash back into  $x$ -th path
  - ▶ If  $\text{op}$  is a read, then return data


## Path ORAM Input




## Worked Example

```
access(0,write,‘‘Zero’’)
```

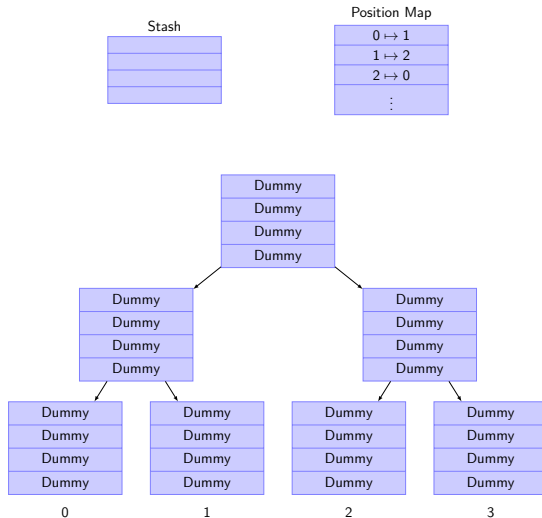
# Example Write: Lookup Position




# Example Write: Remap Block



# Example Write: Read Path




# Example Write: Write Data





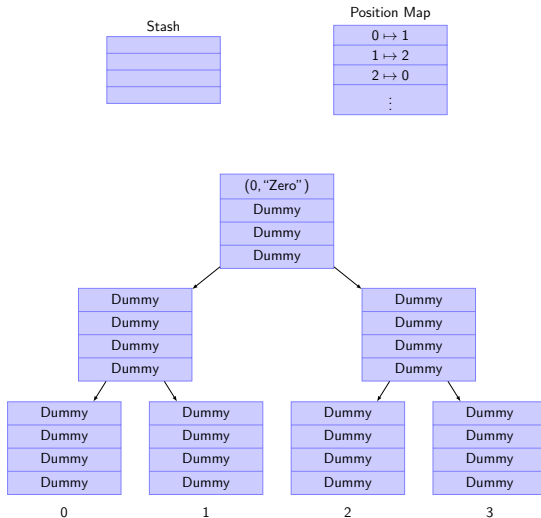
# Example Write: Write Path




## Worked Example

```
access(1,write,‘‘One’’)
```

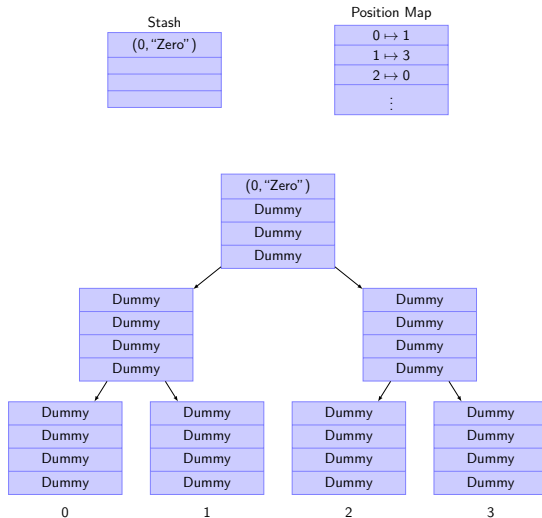
# Example Write: Lookup Position



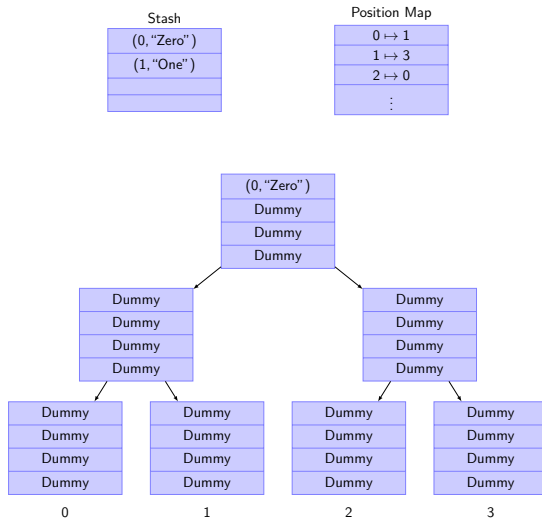
# Example Write: Remap Block



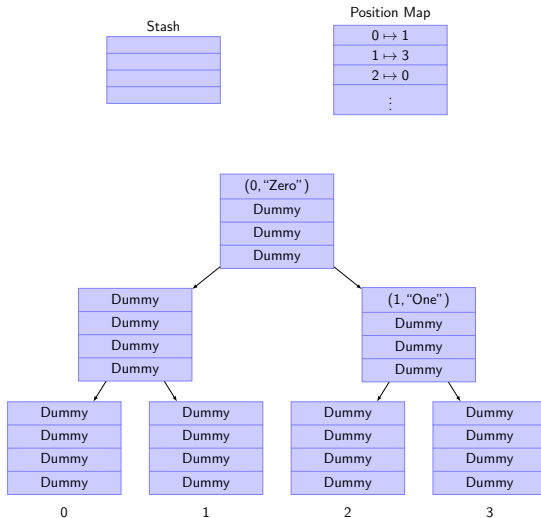
# Example Write: Read Path



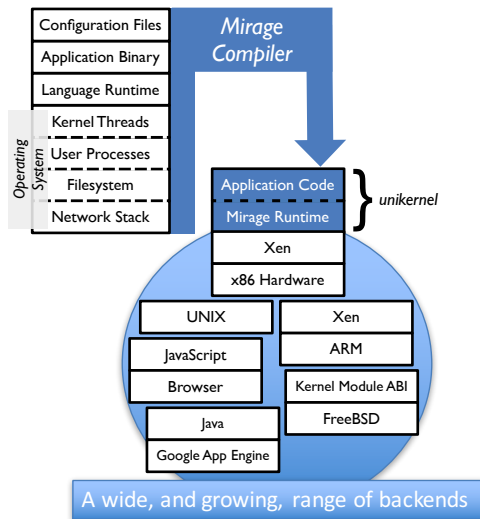
# Example Write: Write Data



# Example Write: Write Path



# MirageOS





# Overview

Motivation

Solution

**Implementation**

Evaluation

Summary

# Basic ORAM


Disk



# Basic ORAM

Encryption


Disk



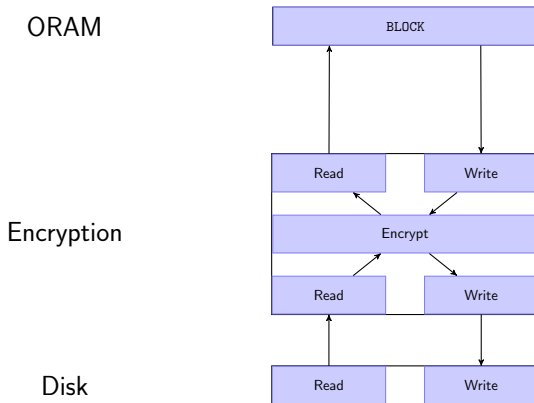
# Basic ORAM

Encryption

Disk




# Basic ORAM




# Basic ORAM

ORAM



Encryption



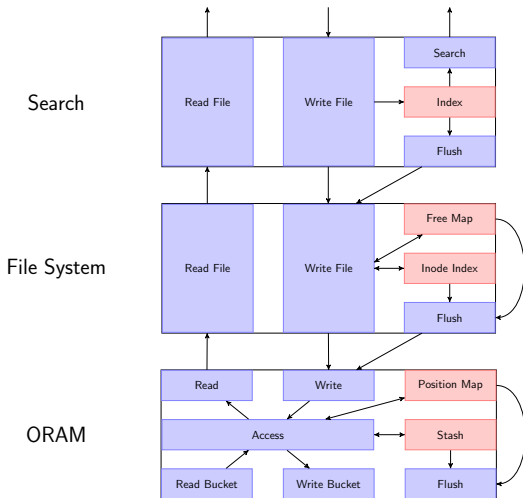
Disk



# Recursive ORAM

- ▶ We want ORAM to be stateless, but writing position map to disk is expensive
- ▶ Recursive ORAM stores the position map of the first ORAM in another ORAM
  - ▶ The second ORAM is smaller than the first
  - ▶ This can be repeated
- ▶ Implemented this using recursive functors

# Search Application





# Overview

Motivation

Solution

Implementation


**Evaluation**

Summary

# Evaluation


- ▶ Explored parameter space
  - ▶ Specifically looked at block size
  - ▶ Increasing block size increased speed
  - ▶ Chose block size of 1MB
- ▶ Measured performance
  - ▶ Compared ORAM with encryption, ORAM without encryption, and control
  - ▶ Showed expected logarithmic overheads
  - ▶ Took  $\approx 1000$ s to transfer 1GB on 4GB ORAM
- ▶ Showed security properties using statistical techniques

# Block Size Results



**Figure:** Plot of the time taken to transfer 80MB of data at varying block sizes and sizes of ORAM. Each line represents one ORAM size,  $N$ , so as block size increases, the time decreases.

# Performance Results



**Figure:** The relationship between size of an ORAM in blocks and the time taken for 1000 operations, plotted for ORAM, encrypted ORAM, and a control block device with no ORAM. We take logs of both axes, because block size was increased in powers of two and we expect a log relationship.


## 1. Autocorrelation plotting

- ▶ Plot the correlation of a sequence with itself for a number of lags
- ▶ For a random sequence noise cancels out to give values close to zero


## 2. Runs testing

- ▶ This counts the number of runs of consecutive values all above or below the median
- ▶ We compare this number to that of a random process

# Autocorrelation Results




(a) Autocorrelation plot of a 200 iteration access pattern



(b) Autocorrelation plot of a 5,000 iteration access pattern

**Figure:** Two autocorrelation plots, with the autocorrelation coefficient on the y-axis and time lag on the x-axis. The dashed black lines represent confidence bands of 95% and 99%. For a random sequence, most of the points should fall within the 95% confidence bound, as they do on both of these plots.

# Runs Test Results



**Figure:** The distribution of the number of runs in 1000 access patterns of length 180. The dashed black lines represent 5% tail cut-offs. 92.2% of values fall within these bounds, implying that the access patterns were created from a random process.

# Overview

Motivation

Solution

Implementation

Evaluation

Summary



# Summary

- ▶ Homomorphic methods of encrypted search can leak information via the side channel of access pattern
- ▶ ORAM provides a solution to this problem
- ▶ My implementation gives the desired security properties while maintaining acceptable performance

Thank You

Questions?

<https://github.com/ruhatch/mirage-oram>